

Problem Set 5

This problem set is due in class on **Thursday, March 28, 1:30pm**.

If you have questions, please post them on the Discussion Board on the Blackboard course website, rather than emailing the course staff. This will allow other students with the same question to see the response and any ensuing discussion.

The dataset for the following problems is posted on the Blackboard course website, under “Course Content → Data sets → ps5_data.mat”. When you load the .mat file, you will find the following variables:

RealWaveform: 10 seconds of real electrode data¹ sampled at 30 kHz. Values are in μV .

InitTwoClusters_1, **InitTwoClusters_2**, **InitThreeClusters_1**, **InitThreeClusters_2**: each is a $31 \times K$ matrix, where the k th column is the initialization of the k th cluster center, μ_k ($k = 1, \dots, K$).

1. (5 points) Data pre-processing

Plot the entire waveform. You will notice a low frequency component in the signal, known as the local field potential (LFP). We are not interested in analyzing the LFP during spike sorting. Remove the LFP using a high-pass filter, stopping frequencies below 250 Hz. You may find the following code useful for designing a simple high-pass filter:

```
x = RealWaveform;
f_0 = 30000; % sampling rate of waveform (Hz)
f_stop = 250; % stop frequency (Hz)
f_Nyquist = f_0/2; % the Nyquist limit
n = length(x);
f_all = linspace(-f_Nyquist,f_Nyquist,n);
desired_response = ones(n,1);
desired_response(abs(f_all)<=f_stop) = 0;
x_filtered = real(fftshift(ifft(fft(x).*fftshift(desired_response))));
```

Plot the filtered waveform. Was the low-frequency component removed?

¹The neural data have been generously provided by the laboratory of Prof. Krishna Shenoy at Stanford University. The data are to be used exclusively for educational purposes in this course.

You can listen to the neurons with this command (only supported on some architectures):

```
sound(x*.97/max(abs(x)),f_0);
```

2. (7 points) Spike detection

Set a threshold, $V_{\text{thresh}} = 250 \mu\text{V}$, and determine all of the times that the signal crosses from below to above the threshold. Plot the threshold as a line across your plot of the high-pass filtered waveform from Problem 1.

Take 1 ms snippets of the waveform beginning 0.3 ms before each threshold crossing. Each snippet should have 31 samples; the tenth sample should be less than V_{thresh} , and the eleventh sample should be greater than V_{thresh} . In a new figure, make a “voltage versus time” plot containing the following:

- all of the threshold-crossing waveform snippets, and
- the threshold as a horizontal line.

Your result should look similar to Figure 3 from Lewicki’s review paper, with the key difference that your plot should show multiple spike shapes (instead of just one). You will also notice a few non-stereotyped traces resulting from either different neurons spiking near the same time or noise exceeding threshold.

3. Clustering with the K-means algorithm

Implement the K-means algorithm in MATLAB, and use it to determine the neuron responsible for each recorded spike.

Treat each snippet as a point $\mathbf{x}_n \in \mathbf{R}^D$ ($n = 1, \dots, N$), where $D = 31$ is the number of samples in each snippet, and N is the number of detected spikes. In this problem, we will assume that there are $K = 2$ neurons contributing spikes to the recorded waveform. Initialize the cluster centers using `InitTwoClusters_1`.

- (13 points) For each cluster ($k = 1, 2$), create a separate “voltage versus time” plot containing the following:
 - the cluster center $\boldsymbol{\mu}_k$ returned by K-means as a red waveform trace (i.e., the prototypical action potential for the k th neuron),
 - all of the waveform snippets assigned to the k th neuron, and
 - the threshold as a horizontal line.
- (10 points) Plot the objective function J versus iteration number (as in Figure 9.2 in *PRML*). How many iterations did it take for K-means to converge?

4. **(5 points)** As discussed in class, K-means guarantees convergence to a local optimum. Thus, it is possible to converge to different local optima with different initializations. Repeat Problem 3 where the cluster centers are initialized using `InitTwoClusters_2`. Is the local optimum found here the same or different as that found in Problem 3?
5. With real data, we don't typically know how many neurons contribute spikes to a given recorded waveform. We will soon discuss in class how to infer the number of neurons from the data, a procedure which typically requires running the clustering algorithm on the same data for different values of K . In this problem, we will assume that there are $K = 3$ (rather than $K = 2$) neurons contributing spikes to the recorded waveform.
 - (a) **(5 points)** Repeat Problem 3 with $K = 3$ and initializing using `InitThreeClusters_1`. How does the local optimum found here differ from that found in Problem 3?
 - (b) **(5 points)** Repeat Problem 3 with $K = 3$ and initializing using `InitThreeClusters_2`. Is the local optimum found here the same or different as that found in part (a)?